



Task2

Helm, Persistent Storage and Ingress



Step1: Containerization

Containerization

Each service moved to a folder

```
./app1:
total 20
drwxrwxr-x  2 master master 4096 Sep  8 14:04 .
drwxrwxrwx 11 root   root  4096 Sep  8 14:06 ./
-rw-rw-r--  1 master master 198 Sep  8 13:00 Dockerfile.app1
-rw-rw-r--  1 master master 454 Sep  8 13:00 log_service_1.py
-rw-rw-r--  1 master master    7 Sep  8 14:04 requirements.txt

./app2:
total 20
drwxrwxr-x  2 master master 4096 Sep  8 14:04 .
drwxrwxrwx 11 root   root  4096 Sep  8 14:06 ./
-rw-rw-r--  1 master master 198 Sep  8 13:00 Dockerfile.app2
-rw-rw-r--  1 master master 454 Sep  8 13:00 log_service_2.py
-rw-rw-r--  1 master master    7 Sep  8 14:04 requirements.txt

./app3:
total 20
drwxrwxr-x  2 master master 4096 Sep  8 14:04 .
drwxrwxrwx 11 root   root  4096 Sep  8 14:06 ./
-rw-rw-r--  1 master master 198 Sep  8 13:00 Dockerfile.app3
-rw-rw-r--  1 master master 454 Sep  8 13:00 log_service_3.py
-rw-rw-r--  1 master master    7 Sep  8 14:04 requirements.txt
```

Build Commands

Then containerized using following docker commands for each:

```
$ docker build -f Dockerfile.app1 -t log-service-1 .
$ docker build -f Dockerfile.app2 -t log-service-2 .
$ docker build -f Dockerfile.app3 -t log-service-3 .
```

Tagging

```
$ docker tag log-service-1 omargabro/app1:v1
$ docker tag log-service-2 omargabro/app2:v1
$ docker tag log-service-3 omargabro/app3:v1
```

Pushing to docker hub

```
$ docker push omargabro/app1:v1
$ docker push omargabro/app2:v1
$ docker push omargabro/app3:v1
```

Docker hub snapshot

Tags	OS	Vulnerabilities	Last pushed	Size
omargabr0/app3	v1	inactive	2 minutes ago	50.09 MB
omargabr0/app2	v1	inactive	2 minutes ago	50.09 MB
omargabr0/app1	v1	inactive	3 minutes ago	50.09 MB

Step2: Helm Deployment

App1 Folder structure

```
task2/
└── app1/
    ├── Chart.yaml
    ├── values.yaml
    ├── templates/
    │   ├── deployment.yaml
    │   ├── service.yaml
    │   ├── ingress.yaml
    │   ├── hpa.yaml
    │   ├── _helpers.tpl
    │   └── NOTES.txt
    └── tests/
        └── test-connection.yaml
    └── charts/
```

App1:

Values.yaml:

```
replicaCount: 1

image:
  repository: omargabr0/app1
  tag: v1
  pullPolicy: IfNotPresent

namespace: env-1

service:
  type: NodePort
```

```

port: 5001
nodePort: 30011

hpa:
  enabled: true
  minReplicas: 1
  maxReplicas: 5
  CPU:
    averageUtilization: 70
  memory:
    enabled: false
    targetMemoryUtilizationPercentage: 80

ingress:
  enabled: false

serviceAccount:
  create: false

```

Deployment.yaml

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: {{ .Release.Name }}
  labels:
    app: {{ .Release.Name }}
    namespace: {{ .Values.namespace }}
spec:
  replicas: {{ .Values.replicaCount }}
  selector:
    matchLabels:
      app: {{ .Release.Name }}
template:
  metadata:
    labels:
      app: {{ .Release.Name }}
spec:
  containers:
    - name: {{ .Release.Name }}
      image: "{{ .Values.image.repository }}:{{ .Values.image.tag }}"
      imagePullPolicy: {{ .Values.image.pullPolicy }}
      ports:
        - containerPort: {{ .Values.service.port }}
  resources:
    requests:
      cpu: "100m"
    limits:

```

```
cpu: "500m"
```

service.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: {{ .Release.Name }}
  labels:
    app: {{ .Release.Name }}
spec:
  type: {{ .Values.service.type }}
  selector:
    app: {{ .Release.Name }}
  ports:
    - protocol: TCP
      port: {{ .Values.service.port }}
      targetPort: {{ .Values.service.port }}
      {{- if eq .Values.service.type "NodePort" -}}
      nodePort: {{ .Values.service.nodePort }}
      {{- end }}
```

hpa.yaml

```
{{- if .Values.hpa.enabled -}}
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: {{ .Release.Name }}
  labels:
    app: {{ .Release.Name }}
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: {{ .Release.Name }}
  minReplicas: {{ .Values.hpa.minReplicas }}
  maxReplicas: {{ .Values.hpa.maxReplicas }}
  metrics:
    - type: Resource
      resource:
        name: cpu
        target:
          type: Utilization
          averageUtilization: {{ .Values.hpa.CPU.averageUtilization }}
```

```

{{- if .Values.hpa.memory.enabled --}}
- type: Resource
  resource:
    name: memory
    target:
      type: Utilization
      averageUtilization: {{ .Values.hpa.targetMemoryUtilizationPercentage }}
{{- end }}

{{- end }}

```

command outputs showing successful deployment.

```

master@master:/usr/new/task2/app1$ helm install app1-release1 ./app1
NAME: app1-release1
LAST DEPLOYED: Tue Sep  9 11:00:20 2025
NAMESPACE: default
STATUS: deployed
REVISION: 1
NOTES:
1. Get the application URL by running these commands:
  export NODE_PORT=$(kubectl get --namespace default -o jsonpath=".spec.ports[0].nodePort" services app1-release1)
  export NODE_IP=$(kubectl get nodes --namespace default -o jsonpath=".items[0].status.addresses[0].address")
  echo http://$NODE_IP:$NODE_PORT
master@master:/usr/new/task2/app1$ helm list
NAME           NAMESPACE   REVISION   UPDATED             STATUS        CHART          APP VERSION
app1-release1  default     1          2025-09-09 11:00:20 459538217 +0000 UTC deployed  app1-0.1.0      1.0
master@master:/usr/new/task2/app1$ helm history app1-release1
REVISION  UPDATED   STATUS   CHART          APP VERSION  DESCRIPTION
1         Tue Sep  9 11:00:20 2025  deployed  app1-0.1.0  1.0          Install complete
master@master:/usr/new/task2/app1$ 

```

App2 Folder structure

```

task2/
└── app2/
    ├── Chart.yaml
    ├── values-env1.yaml
    ├── values-env2.yaml
    └── templates/
        ├── deployment.yaml
        ├── service.yaml
        ├── _helpers.tpl
        ├── NOTES.txt
        └── tests/
            └── test-connection.yaml
    └── charts/

```

Values-env1.yaml

```

replicaCount: 1

image:
  repository: omargabro/app2
  tag: v1

```

```
pullPolicy: IfNotPresent

namespace: env-1

service:
  type: NodePort
  port: 5002
  nodePort: 30012

hpa:
  enabled: true
  minReplicas: 1
  maxReplicas: 5
  CPU:
    averageUtilization: 70
  memory:
    enabled: false
    targetMemoryUtilizationPercentage: 80

ingress:
  enabled: false

serviceAccount:
  create: false
```

Values-env2.yaml

```
replicaCount: 1

image:
  repository: omargabr0/app2
  tag: v1
  pullPolicy: IfNotPresent

namespace: env-2

service:
  type: NodePort
  port: 5002
  nodePort: 30013

hpa:
  enabled: true
  minReplicas: 1
  maxReplicas: 5
  CPU:
    averageUtilization: 70
  memory:
    enabled: false
```

```

targetMemoryUtilizationPercentage: 80

ingress:
  enabled: false

serviceAccount:
  create: false

```

rest of templet files as deployment.yaml, service.yaml, hpa.yaml are the same as app1, only changes where necessary to be made are in the values.yaml.

Command to run using specified values files:

```
$ helm install app2-env1 ./app2 -f /usr/new/task2/app2/app2/values-env1.yaml
```

```
$ helm install app2-env2 ./app2 -f /usr/new/task2/app2/app2/values-env2.yaml
```

command outputs showing successful deployment.

```

NOTES:
1. Get the application URL by running these commands:
  export NODE_PORT=$(kubectl get --namespace default -o jsonpath=".spec.ports[0].nodePort" services app2-env1)
  export NODE_IP=$(kubectl get nodes --namespace default -o jsonpath=".items[0].status.addresses[0].address")
  echo http://$NODE_IP:$NODE_PORT
master@master:/usr/new/task2/app2$ kubectl get pods -n env-2
  NAME        READY   STATUS    RESTARTS   AGE
  app2-env2-bfcb4fbcb-vm4g4   1/1     Running   0          67s
master@master:/usr/new/task2/app2$ █

master@master:/usr/new/task2/app2$ helm status app2-env1
NAME: app2-env1
LAST DEPLOYED: Tue Sep  9 12:38:55 2025
NAMESPACE: default
STATUS: deployed
REVISION: 1
NOTES:
1. Get the application URL by running these commands:
  export NODE_PORT=$(kubectl get --namespace default -o jsonpath=".spec.ports[0].nodePort" services app2-env1)
  export NODE_IP=$(kubectl get nodes --namespace default -o jsonpath=".items[0].status.addresses[0].address")
  echo http://$NODE_IP:$NODE_PORT
master@master:/usr/new/task2/app2$ helm list
  NAME      NAMESPACE   REVISON   UPDATED           STATUS   CHART      APP VERSION
  app1-release1  default      1       2025-09-09 11:00:20.459538217 +0000 UTC deployed  app1-0.1.0   1.0
  app2-env1      default      1       2025-09-09 12:38:55.872008943 +0000 UTC deployed  app2-0.1.0   1.16.0
master@master:/usr/new/task2/app2$ kubectl get pods -n env2

```

App3: Folder structure

Similar to app1

```

task2/
└── app3/
    ├── Chart.yaml
    ├── values.yaml
    └── templates/
        ├── deployment.yaml
        └── service.yaml

```

```
|   ├── ingress.yaml
|   ├── hpa.yaml
|   ├── _helpers.tpl
|   ├── NOTES.txt
|   └── tests/
|       └── test-connection.yaml
└── charts/
```

Values.yaml

```
replicaCount: 1

image:
  repository: omargabro/app3
  tag: v1
  pullPolicy: IfNotPresent

namespace: env-3

service:
  type: NodePort
  port: 5003
  nodePort: 30013

hpa:
  enabled: true
  minReplicas: 1
  maxReplicas: 5
  CPU:
    averageUtilization: 70
  memory:
    enabled: false
    targetMemoryUtilizationPercentage: 80

ingress:
  enabled: false

serviceAccount:
  create: false
```

command outputs showing successful deployment.

```

release app3-env2 unstalled
master@master:/usr/new/task2/app3$ helm install app3-env2 ./app3
NAME: app3-env2
LAST DEPLOYED: Tue Sep  9 14:16:54 2025
NAMESPACE: default
STATUS: deployed
REVISION: 1
NOTES:
1. Get the application URL by running these commands:
  export NODE_PORT=$(kubectl get --namespace default -o jsonpath=".spec.ports[0].nodePort" services app3-env2)
  export NODE_IP=$(kubectl get nodes --namespace default -o jsonpath=".items[0].status.addresses[0].address")
  echo http://$NODE_IP:$NODE_PORT
master@master:/usr/new/task2/app3$ kubectl get pods -n env2
No resources found in env2 namespace.
master@master:/usr/new/task2/app3$ kubectl get pods -n env-2
NAME          READY   STATUS    RESTARTS   AGE
app2-env2-bfcf4fbcb-vm4g4   1/1     Running   0          91m
app3-env2-54ddf68b44-5r7r6   1/1     Running   0          17s
master@master:/usr/new/task2/app3$ kubectl get svc -n env-2
No resources found in env-2 namespace.
master@master:/usr/new/task2/app3$ kubectl get svc
NAME        TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
app1-release1  NodePort  10.110.190.138  <none>        5001:30011/TCP  3h17m
app2-env1     NodePort  10.111.220.101   <none>        5002:30012/TCP  98m
app2-env2     NodePort  10.103.237.36   <none>        5002:30013/TCP  91m
app3-env2     NodePort  10.100.244.87   <none>        5003:30014/TCP  39s
kubernetes    ClusterIP 10.96.0.1       <none>        443/TCP      5d15h
ner-svc        NodePort  10.96.186.221   <none>        80:30002/TCP  4d1h
proxy-svc      NodePort  10.106.12.174   <none>        80:30003/TCP  4d1h
sentiment-svc  NodePort  10.111.180.255  <none>        80:30001/TCP  4d1h
master@master:/usr/new/task2/app3$ █

```

Step 3: Persistent Storage

Creating pv-pvc file in app/template/pv-pvc.yaml so that it will be implemented to helm chart.

App1

Adding pv-pvc:

This will mount the volume to /logs in the container filesystem

And the host store the pv data /usr/new/task2/app1/logs

Values.yaml

```

pv:
  enabled: true
  accessMode: ReadWriteMany
  hostPath: /usr/new/task2/app1/logs
  mountPath: /logs
  persistentVolumeReclaimPolicy: Delete
  storage: 500Mi

```

Deployment.yaml

```

{{- if .Values.pv.enabled }}
  volumeMounts:
    - name: log-volume
      mountPath: {{ .Values.pv.mountPath }}
  {{- end }}
{{- if .Values.pv.enabled }}
  volumes:

```

```

- name: log-volume
  persistentVolumeClaim:
    claimName: {{ .Release.Name }}-pvc
{{- end }}

```

pv-pvc.yaml

```

{{- if .Values.pv.enabled --}}
---
apiVersion: v1
kind: PersistentVolume
metadata:
  name: {{ .Release.Name }}-pv
spec:
  capacity:
    storage: {{ .Values.pv.storage }}
  accessModes:
    - {{ .Values.pv.accessMode }}
  hostPath:
    path: {{ .Values.pv.hostPath }}

---
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: {{ .Release.Name }}-pvc
spec:
  accessModes:
    - {{ .Values.pv.accessMode }}
  resources:
    requests:
      storage: {{ .Values.pv.storage }}
{{- end --}}

```

App2 in env-1 and env-2

Same yaml files, but the difference in values file.

Values-env1.yaml

```

pv:
  enabled: true
  accessMode: ReadWriteMany
  hostPath: /usr/new/task2/app2/logs-env1
  mountPath: /logs

```

```
persistentVolumeReclaimPolicy: Delete
storage: 500Mi
```

Values-env1.yaml

```
pv:
  enabled: true
  accessMode: ReadWriteMany
  hostPath: /usr/new/task2/app2/logs-env2
  mountPath: /logs
  persistentVolumeReclaimPolicy: Delete
  storage: 500Mi
```

App3

Creating pv-pvc file in app/template/pv-pvc.yaml so that it will be implemented to helm chart.

Values.yaml

```
pv:
  enabled: true
  accessMode: ReadWriteMany
  hostPath: /usr/new/task2/app3/logs
  mountPath: /logs
  persistentVolumeReclaimPolicy: Delete
  storage: 500Mi
```

Final helm application and pv- pvc snapshots:

```

master@master:~$ helm list
NAME      NAMESPACE   REVISION    UPDATED             STATUS        CHART          APP VERSION
app1-env1  default     1          2025-09-10 12:36:37.437597322 +0000 UTC deployed  app1-0.1.0    1.0
app2-env1  default     1          2025-09-10 12:24:36.055245865 +0000 UTC deployed  app3-0.1.0    1.0
app2-env2  default     1          2025-09-10 12:24:42.123441351 +0000 UTC deployed  app3-0.1.0    1.0
app3-env3  default     1          2025-09-10 12:33:09.911245765 +0000 UTC deployed  app3-0.1.0    1.0
master@master:~$ kubectl get pv && kubectl get pvc
NAME           CAPACITY  ACCESS MODES  RECLAIM POLICY  STATUS   CLAIM          STORAGECLASS  VOLUMEATTRIBUTESCLASS  REASON  AGE
app1-env1-pv   500Mi     RWX         Delete        Bound    default/app1-env1-pvc  <unset>       5h1m
app2-env1-pv   500Mi     RWX         Delete        Bound    default/app2-env1-pvc  <unset>       5h13m
app2-env2-pv   500Mi     RWX         Delete        Bound    default/app2-env2-pvc  <unset>       5h13m
app3-env3-pv   500Mi     RWX         Delete        Bound    default/app3-env3-pvc  <unset>       5h5m
NAME           STATUS     VOLUME      CAPACITY  ACCESS MODES  STORAGECLASS  VOLUMEATTRIBUTESCLASS  AGE
app1-env1-pvc  Bound     app1-env1-pv  500Mi     RWX          <unset>       5h1m
app2-env1-pvc  Bound     app2-env1-pv  500Mi     RWX          <unset>       5h13m
app2-env2-pvc  Bound     app2-env2-pv  500Mi     RWX          <unset>       5h13m
app3-env3-pvc  Bound     app3-env3-pv  500Mi     RWX          <unset>       5h5m
master@master:~$ kubectl get svc -n env-1 env-2
Error from server (NotFound): services "env-2" not found
master@master:~$ kubectl get svc -n env-1 -n env-2 -n env-3
NAME      TYPE    CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
app3-env3  NodePort  10.100.1.90  <none>          5003:30014/TCP  5h5m
master@master:~$ kubectl get svc -n env-1
NAME      TYPE    CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
app1-env1  NodePort  10.106.40.162  <none>          5001:30011/TCP  5h2m
app2-env1  NodePort  10.111.250.191  <none>          5002:30012/TCP  5h14m
master@master:~$ kubectl get svc -n env-2
NAME      TYPE    CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
app2-env2  NodePort  10.111.59.123  <none>          5002:30013/TCP  5h14m
master@master:~$ kubectl get svc -n env-3
NAME      TYPE    CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
app3-env3  NodePort  10.100.1.90  <none>          5003:30014/TCP  5h6m

```

This snapshot shows:

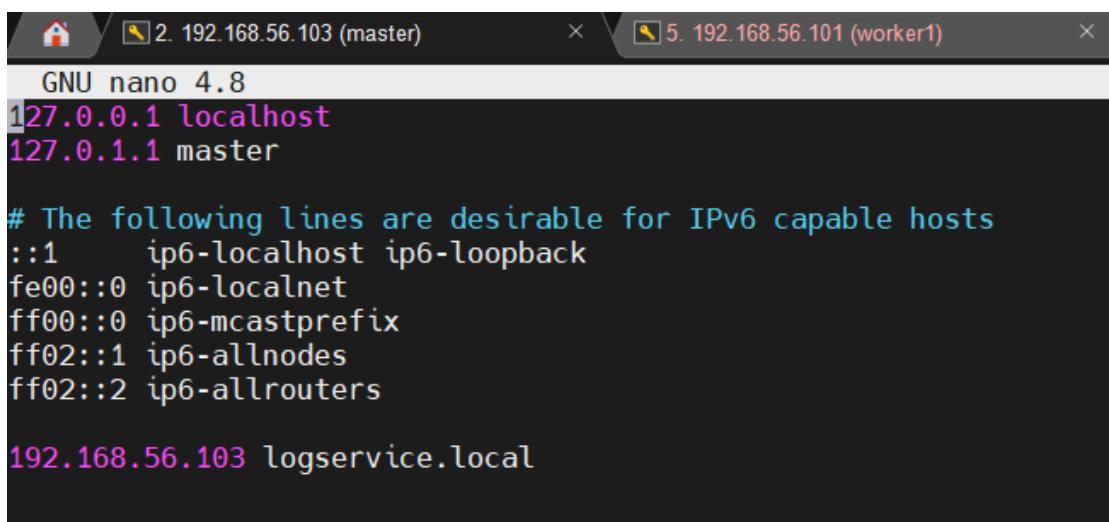
- all helm application
- pv and pvc created successfully
- services for each application

Step 4: Ingress

DNS configuration

This maps the master IP to domain so that we can use in ingress resource file

\$ sudo nano /etc/hosts



The screenshot shows two terminal windows side-by-side. Both windows are titled 'GNU nano 4.8'. The left window, labeled '2. 192.168.56.103 (master)', contains the following hosts entries:

```

127.0.0.1 localhost
127.0.1.1 master
# The following lines are desirable for IPv6 capable hosts
::1      ip6-localhost ip6-loopback
fe00::0  ip6-localnet
ff00::0  ip6-mcastprefix
ff02::1  ip6-allnodes
ff02::2  ip6-allrouters
192.168.56.103 logservice.local

```

The right window, labeled '5. 192.168.56.101 (worker1)', contains the same hosts entries.

```
master@master:/usr/new/task2/ingress$ kubectl apply -f ingress.yaml
ingress.networking.k8s.io/my-apps-ingress created
master@master:/usr/new/task2/ingress$ kubectl get validatingwebhookconfiguration ingress-nginx-admission -o yaml | grep failurePolicy
  {"apiVersion":"admissionregistration.k8s.io/v1","kind":"ValidatingWebhookConfiguration","metadata":{},"labels":{"app.kubernetes.io/component":"admission-webhook","app.kubernetes.io/instance":"ingress-nginx","app.kubernetes.io/name":"ingress-nginx","app.kubernetes.io/part-of":"ingress-nginx","app.kubernetes.io/version":"1.13.2"},"name":"ingress-nginx-admission","webhooks":[{"admissionReviewVersions":["v1"],"clientConfig":{"service":{"name":"ingress-nginx-controller-admission","namespace":"ingress-nginx","path":"/networking/v1/ingresses","port":443}),"failurePolicy":"Fail","matchPolicy":"Equivalent","name":"validate.nginx.ingress.kubernetes.io","rules":[{"apiGroups":["networking.k8s.io"],"apiVersions":["v1"],"operations":["CREATE","UPDATE"],"resources":["ingresses"]}], "sideEffects":"None"}]},{ "failurePolicy": "Ignore" }
```

Applying ingress

NAMESPACE	NAME	CLASS	HOSTS	ADDRESS	PORTS	AGE
default	my-apps-ingress	nginx	logservice.loc		80	17m

1. Fix service

NodePort → Cluster IP

Exposing services using NodePort will break the routing as ingress forwards traffic inside the cluster directly to service endpoints.

2. Using ExternalName service

- create aliases of services in each namespace inside the same namespace as the Ingress so that ingress can see it from its namespace
- aliases points to the services in other namespaces like env-1, and ingress routes to the aliases

```
apiVersion: v1
kind: Service
metadata:
  name: app1-env1-alias
  namespace: default
spec:
  type: ExternalName
  externalName: app1-env1.env-1.svc.cluster.local
  ports:
    - port: 5001
      protocol: TCP
```

3. ingress.yaml

```

apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: my-apps-ingress
  namespace: default
  annotations:
    kubernetes.io/ingress.class: nginx
    nginx.ingress.kubernetes.io/rewrite-target: /log
spec:
  rules:
  - host: logservice.local
    http:
      paths:
      - path: /app1
        pathType: Prefix
        backend:
          service:
            name: app1-env1-alias
            port:
              number: 5001
      - path: /app2/env1
        pathType: Prefix
        backend:
          service:
            name: app2-env1-alias
            port:
              number: 5002
      - path: /app2/env2
        pathType: Prefix
        backend:
          service:
            name: app2-env2-alias
            port:
              number: 5002
      - path: /app3
        pathType: Prefix
        backend:
          service:
            name: app3-env3-alias
            port:
              number: 5003

```

- using rewrite option to forward the request to the end point **/log**

4. testing

```
master@master:/usr/new/task2/ingress$ curl -X POST http://logservice.local:30343/app1/log -H "Content-Type: application/json" {"status":"logged"}
master@master:/usr/new/task2/ingress$ curl -X POST http://logservice.local:30343/app2A/log -H "Content-Type: application/json" {"status":"logged"}
master@master:/usr/new/task2/ingress$ curl -X POST http://logservice.local:30343/app2B/log -H "Content-Type: application/json" {"status":"logged"}
master@master:/usr/new/task2/ingress$ curl -X POST http://logservice.local:30343/app3/log -H "Content-Type: application/json" {"status":"logged"}  
[{"status": "logged"}]
```

- from the results of snapshots, endpoints are reachable after using ingress
- logservice.local resolved to the master IP which is the ingress host
- ingress resolve paths to app1.env-1.svc.cluster.local

Step 5: Automation scripts

Deploy.sh

```
#!/bin/bash
set -e

VERSION=$(date +%Y%m%d%H%M)      # timestamp as tag

show_usage()
{
    echo "Deployment script V 1.0"
    echo "Usage: $0 [[--image=app1 or app2 or app3 ][--username= user personal docker hub]"  
] [ --appsdirectory=directory of build] [ --release= helm release name] [ --namespace = env1]]"
}

show_help()
{
    show_usage
    echo "  where"
    echo "  --image: the app image name e.g app1 , app2 , app3"
    echo "  --username: Docker hub and the name tag of image"
    echo "  --appsdirectory: Apps directory that have each app Docker file e.g ./app1/Dockerfile"
    echo "  --release: Release name of helm chart "
    echo "  --namespace: namespace of app e.g env1 , env2 "

}
DEFAULTS() {
    DOCKER_USER=omargabro
    DOCKER_DIR=/usr/new/task2/build-dir
    ENV=default
```

```

}

parse_args() {
    DEFAULTS
for i in "$@"; do
    case $i in

        --username=*) DOCKER_USER=${i#*=} ;;
        --appsdirectory=*) DOCKER_DIR=${i#*=} ;;
        --release=*) RELEASE=${i#*=} ;;
        --image=*) IMAGE_NAME=${i#*=} ;;
        --namespace=*) ENV=${i#*=} ;;
        --help)      show_help; exit 0 ;;
        *)          show_usage; exit 1 ;;
    esac
done
}

#login check
login(){
if docker info 2>/dev/null | grep -q "Username: $DOCKER_USER"; then
    echo " Already logged in to Docker Hub"
else
    echo " Not logged in. Login to docker hub first"
    exit 0
fi
}

# Build image
build(){
echo ">>> Building Docker image..."
docker build -f $DOCKER_DIR/$IMAGE_NAME/Dockerfile.$IMAGE_NAME \
-t $DOCKER_USER/$IMAGE_NAME:$VERSION \
$DOCKER_DIR/$IMAGE_NAME
}

# Push image
push(){
echo ">>> Pushing Docker image to Docker Hub..."
docker push $DOCKER_USER/$IMAGE_NAME:$VERSION
}

# Update Kubernetes deployment
update() {
    if [ "$IMAGE_NAME" = "app2" ]; then
        helm upgrade $RELEASE ../$IMAGE_NAME/$IMAGE_NAME -f values-$ENV.yaml --set
        image.tag=$VERSION
    else
        echo ">>> Updating Kubernetes deployment Using Helm charts..."
    fi
}

```

```

    helm upgrade $RELEASE ./${IMAGE_NAME}/$IMAGE_NAME --set image.tag=$VERSION
fi
}

main (){
    parse_args "$@"
    login
    build
    push
    update

echo ">>> Deployment successful with image ${IMAGE_NAME}:${VERSION}"
echo ">>> Helm chart $RELEASE updated with image ${IMAGE_NAME}:${VERSION}"
}

main "$@"

```

- **Automation:** Builds, pushes, and deploys Docker images automatically.
- **Versioning:** Uses timestamp tags to ensure each image is unique.
- **Argument Handling:** Supports flexible inputs (image, user, release, namespace).
- **Defaults:** Provides fallback values if no arguments are given.
- **Login Check:** Ensures Docker Hub authentication before pushing.
- **Build & Push:** Creates Docker images from app Dockerfiles and pushes them to Docker Hub.
- **Deployment Update:** Uses Helm to update Kubernetes releases with the new image.
- **Reliability:** Script exits immediately on errors to avoid broken deployments.

Test 1: app1-env1

Rollout.sh

```

#!/bin/bash
set -e

show_usage() {
    echo "Rollback Script V1.0"
    echo "Usage: $0 [--release=helm_release] [--revision=number]"
}

```

```

}

DEFAULTS() {
    NAMESPACE=default
}

parse_args() {
    DEFAULTS
    for i in "$@"; do
        case $i in
            --release=*) RELEASE=${i#*=} ;;
            --namespace=*) NAMESPACE=${i#*=} ;;
            --revision=*) REVISION=${i#*=} ;;
            --help) show_usage; exit 0 ;;
            *) show_usage; exit 1 ;;
        esac
    done
}

rollback() {
    if [ -z "$RELEASE" ]; then
        echo "Error: --release is required"
        show_usage
        exit 1
    fi

    if [ -n "$REVISION" ]; then
        echo ">>> Rolling back Helm release $RELEASE in namespace $NAMESPACE to revision $REVISION..."
        helm rollback $RELEASE $REVISION -n $NAMESPACE
    else
        echo ">>> Rolling back Helm release $RELEASE in namespace $NAMESPACE to previous revision..."
        helm rollback $RELEASE -n $NAMESPACE
    fi
}

main() {
    parse_args "$@"
    rollback
    echo ">>> Rollback successful for release $RELEASE"
}

main "$@"

```